

In [1]: `from google.colab import drive`

Go to this URL in a browser: https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3aietf%3awg%3aoauth%3a2.0%3aob&response_type=code&scope=email%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdocs.test%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive.photos.readonly%20https%3a%2f%2fwww.googleapis.com%2fauth%2fpeopleapi.readonly (https://accounts.google.com/o/oauth2/auth?client_id=947318989803-6bn6qk8qdgf4n4g3pfee6491hc0brc4i.apps.googleusercontent.com&redirect_uri=urn%3aietf%3awg%3aoauth%3a2.0%3aob&response_type=code&scope=email%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdocs.test%20https%3a%2f%2fwww.googleapis.com%2fauth%2fdrive.photos.readonly%20https%3a%2f%2fwww.googleapis.com%2fauth%2fpeopleapi.readonly)

Enter your authorization code:

.....

Mounted at /content/drive

In [2]: `!cp /content/drive/My Drive/21. Transfer Learning/rvl-cdip.rar`

In [3]: `!python3 /content/drive/My Drive/21. Transfer Learning/rvl-cdip.py`

Streaming output truncated to the last 5000 lines.

Creating	data_final/imagesy/y/w/v/ywv07e00	OK	
Extracting	data_final/imagesy/y/w/v/ywv07e00/2031300697.tif	95	OK
Creating	data_final/imagesy/y/w/v/ywv14c00	OK	
Extracting	data_final/imagesy/y/w/v/ywv14c00/2080388930b.tif	95	OK
Creating	data_final/imagesy/y/w/v/ywv15a00	OK	
Extracting	data_final/imagesy/y/w/v/ywv15a00/528410291+-0291.tif	95	OK
Creating	data_final/imagesy/y/w/v/ywv90a00	OK	
Extracting	data_final/imagesy/y/w/v/ywv90a00/0060014591.tif	95	OK
Creating	data_final/imagesy/y/w/v/ywv90c00	OK	
Extracting	data_final/imagesy/y/w/v/ywv90c00/13527205.tif	95	OK
Creating	data_final/imagesy/y/w/v/ywv96c00	OK	
Extracting	data_final/imagesy/y/w/v/ywv96c00/CTRCONTRACTS015375-5.tif	95	OK
Creating	data_final/imagesy/y/w/w	OK	
Creating	data_final/imagesy/y/w/w/yww13e00	OK	
Extracting	data_final/imagesy/y/w/w/yww13e00/2061012465.tif	95	OK
Creating	data_final/imagesy/y/w/w/yww44a00	OK	
Extracting	data_final/imagesy/y/w/w/yww44a00/93807744_7748.tif	95	OK
Creating	data_final/imagesy/y/w/w/yww46d00	OK	
Extracting	data_final/imagesy/y/w/w/yww46d00/50506015_6015.tif	95	OK

1. Download all the data in this folder <https://drive.google.com/open?id=1Z4TyI7FcFVEx8qdl4jO9qxvxaqLSqoEu>. it contains two file both images and labels. The label file list the images and their categories in the following format:

path/to/the/image.tif,category

where the categories are numbered 0 to 15, in the following order:

- 0 letter
- 1 form
- 2 email
- 3 handwritten
- 4 advertisement
- 5 scientific report
- 6 scientific publication
- 7 specification
- 8 file folder
- 9 news article
- 10 budget
- 11 invoice
- 12 presentation
- 13 questionnaire
- 14 resume
- 15 memo

2. On this image data, you have to train 3 types of models as given below. You have to split the data into Train and Validation data.

3. Try not to load all the images into memory, use the generators that we have given the reference notebooks to load the batch of images only during the train data.

or you can use this method also

<https://medium.com/@vijayabhaskar96/tutorial-on-keras-imagedatagenerator-with-flow-from-dataframe-8bd5776e45c1> (<https://medium.com/@vijayabhaskar96/tutorial-on-keras-imagedatagenerator-with-flow-from-dataframe-8bd5776e45c1>)

<https://medium.com/@vijayabhaskar96/tutorial-on-keras-flow-from-dataframe-1fd4493d237c> (<https://medium.com/@vijayabhaskar96/tutorial-on-keras-flow-from-dataframe-1fd4493d237c>)

4. You are free to choose Learning rate, optimizer, loss function, image augmentation, any hyperparameters. but you have to use the same architecture what

Start of experiment

```
In [4]: import pandas as pd

data = pd.read_csv('labels_final.csv') #reading the csv file

from sklearn.model_selection import train_test_split

train_path, validation_path, train_label, validation_label = train_test_split(data['p

len(train_path)
```

Out[4]: 38400

```
In [5]: import os
labels_dict = { 0 : 'letter',1: 'form',2: 'email',3 : 'handwritten',4 : 'advertisement'
               5 : 'scientific report',6 : 'scientific publication',7 : 'specificatio
               9 : 'news article', 10 : 'budget', 11 : 'invoice',12 : 'presentation'
               13 : 'questionnaire', 14 : 'resume', 15 : 'memo'}

labels_dict.values()

for subfolder_name in list(labels_dict.values()):
    os.makedirs(os.path.join('train_images', subfolder_name), exist_ok=True)
```

```
In [6]: len(labels_dict)

os.listdir('train_images')
#https://thispointer.com/python-how-to-copy-files-from-one-location-to-another-using-
import shutil
from tqdm import tqdm
for file, label in tqdm(zip(train_path, train_label)):
    shutil.copy('data_final/'+file, 'train_images/'+labels_dict[label]+'/')

for subfolder_name in list(labels_dict.values()):
    os.makedirs(os.path.join('validation_images', subfolder_name))
for file, label in tqdm(zip(validation_path, validation_label)):
    shutil.copy('data_final/'+file, 'validation_images/'+labels_dict[label]+'/')

dir_path= 'train_images'
for i in os.listdir(dir_path):
    print("No of Images in ",i," category is ",len(os.listdir(os.path.join(dir_path,i
dir_path= 'validation_images'
for i in os.listdir(dir_path):
    print("No of Images in ",i," category is ",len(os.listdir(os.path.join(dir_path,i

import tensorflow as tf
dir_path= 'train_images'
ImageFlow = tf.keras.preprocessing.image.ImageDataGenerator()
ImageGenerator_train = ImageFlow.flow_from_directory(dir_path,target_size=(224,224),s
dir_path='validation_images'
ImageGenerator_validation = ImageFlow.flow_from_directory(dir_path,target_size=(224,2

38400it [02:00, 319.56it/s]
9600it [00:35, 268.58it/s]
```

```

No of Images in file folder category is 2351
No of Images in memo category is 2401
No of Images in scientific publication category is 2373
No of Images in handwritten category is 2413
No of Images in invoice category is 2397
No of Images in specification category is 2391
No of Images in scientific report category is 2379
No of Images in letter category is 2461
No of Images in budget category is 2422
No of Images in form category is 2407
No of Images in advertisement category is 2393
No of Images in news article category is 2392
No of Images in resume category is 2415
No of Images in email category is 2379
No of Images in presentation category is 2428
No of Images in questionnaire category is 2395
No of Images in file folder category is 652
No of Images in memo category is 595
No of Images in scientific publication category is 612
No of Images in handwritten category is 592
No of Images in invoice category is 595
No of Images in specification category is 609
No of Images in scientific report category is 620
No of Images in letter category is 554
No of Images in budget category is 580
No of Images in form category is 587
No of Images in advertisement category is 601
No of Images in news article category is 609
No of Images in resume category is 590
No of Images in email category is 614

```

```
In [7]: log_dir="logs/fit/model-1"
```

```
In [8]:
```

Model-1

1. Use [VGG-16 \(https://www.tensorflow.org/api_docs/python/tf/keras/applications/VGG16\)](https://www.tensorflow.org/api_docs/python/tf/keras/applications/VGG16) pretrained network without Fully Connected layers and initialize all the weights with Imagenet trained weights.
2. After VGG-16 network without FC layers, add a new Conv block (1 Conv layer and 1 Maxpooling), 2 FC layers and a output layer to classify 16 classes. You are free to choose any hyperparameters/parameters of conv block, FC layers, output layer.
3. Final architecture will be **INPUT --> VGG-16 without Top layers(FC) --> Conv Layer --> Maxpool Layer --> 2 FC layers --> Output Layer**
4. Train only new Conv block, FC layers, output layer. Don't train the VGG-16 network.

```
In [9]: '''
        This code is refenced from the keras offical documentation page for creating Image
        '''

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D
from tensorflow.keras.layers import Activation, Dropout, Flatten, Dense
```

```

import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator
batch_size = 64
# this is the augmentation configuration we will use for training
train_datagen = ImageDataGenerator(
    rescale=1./255,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True)

# this is the augmentation configuration we will use for testing:
# only rescaling
test_datagen = ImageDataGenerator(rescale=1./255)

tf.keras.backend.clear_session()
!rm -rf logs/
#import keras
#keras.backend.clear_session()
from tensorflow.keras.applications.vgg16 import VGG16
#import tensorflow as tf
from tensorflow.keras.models import Model
from tensorflow.keras.optimizers import Adam
#nb_epochs = 2
batch_size = 64
nb_classes = len(labels_dict)

vgg16_model = VGG16(input_shape=(224,224,3),weights = 'imagenet', include_top = False)
x = vgg16_model.output
x = Conv2D(64, (3, 3), padding="same", kernel_initializer=tf.keras.initializers.he_normal)(x)
x = MaxPooling2D(pool_size=(2, 2))(x)
x = Flatten()(x)
x = Dense(64, kernel_initializer=tf.keras.initializers.he_normal(seed=None), activation='relu')(x)
x = Dense(20, kernel_initializer=tf.keras.initializers.he_normal(seed=None), activation='relu')(x)
predictions = Dense(16, activation = 'softmax')(x)
model = Model(inputs = vgg16_model.input, outputs = predictions)
for layers in vgg16_model.layers:
    layers.trainable = False
model.summary()

tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir,histogram_freq=10)

model.compile(optimizer=Adam(learning_rate=0.001),loss='categorical_crossentropy',metrics=['accuracy'])
history = model.fit_generator(
    ImageGenerator_train,
    steps_per_epoch=38400//32,
    epochs=5,
    validation_data=ImageGenerator_validation,validation_steps= 9600//32, callback=
    tensorboard_callback)
Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5 (https://storage.googleapis.com/tensorflow/keras-applications/vgg16/vgg16_weights_tf_dim_ordering_tf_kernels_notop.h5)
58892288/58889256 [=====] - 1s 0us/step
Model: "model"

```

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0

block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
conv2d (Conv2D)	(None, 7, 7, 64)	294976
max_pooling2d (MaxPooling2D)	(None, 3, 3, 64)	0
flatten (Flatten)	(None, 576)	0
dense (Dense)	(None, 64)	36928
dense_1 (Dense)	(None, 20)	1300
dense_2 (Dense)	(None, 16)	336

=====

Total params: 15,048,228
 Trainable params: 333,540
 Non-trainable params: 14,714,688

WARNING:tensorflow:From <ipython-input-9-396c200b8a3e>:54: Model.fit_generator (from tensorflow.python.keras.engine.training) is deprecated and will be removed in a future version.


Instructions for updating:
 Please use Model.fit, which supports generators.

Epoch 1/5
 1200/1200 [=====] - 296s 247ms/step - loss: 2.7930 - accuracy: 0.1569 - val_loss: 2.3942 - val_accuracy: 0.3074
 Epoch 2/5
 1200/1200 [=====] - 286s 238ms/step - loss: 1.8877 - accuracy: 0.4765 - val_loss: 1.6503 - val_accuracy: 0.5476
 Epoch 3/5
 1200/1200 [=====] - 269s 224ms/step - loss: 1.4253 - accuracy: 0.6131 - val_loss: 1.4141 - val_accuracy: 0.6157
 Epoch 4/5
 1200/1200 [=====] - 258s 215ms/step - loss: 1.2588 - accuracy: 0.6687 - val_loss: 1.3624 - val_accuracy: 0.6515
 Epoch 5/5

```
1200/1200 [=====] - 251s 209ms/step - loss: 1.1661 - accu
racy: 0.7042 - val loss: 1.3809 - val accuracy: 0.6596
```

```
In [10]: from tensorflow.keras.utils import plot_model
```

```
Out[10]:
```

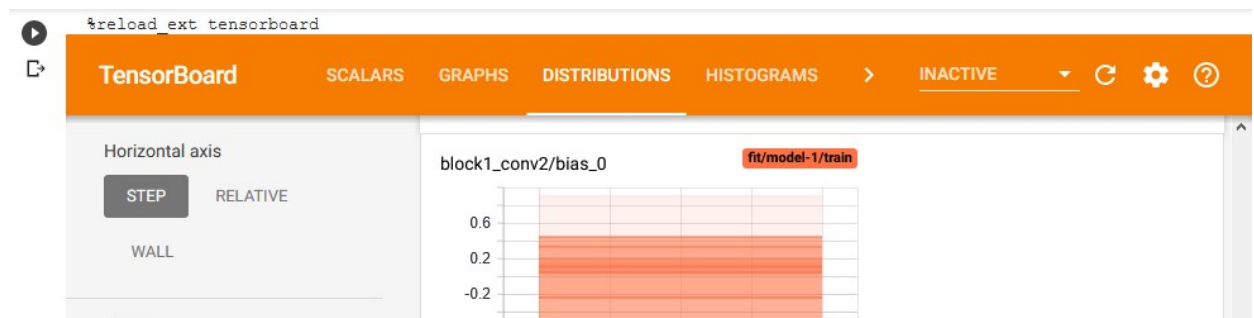



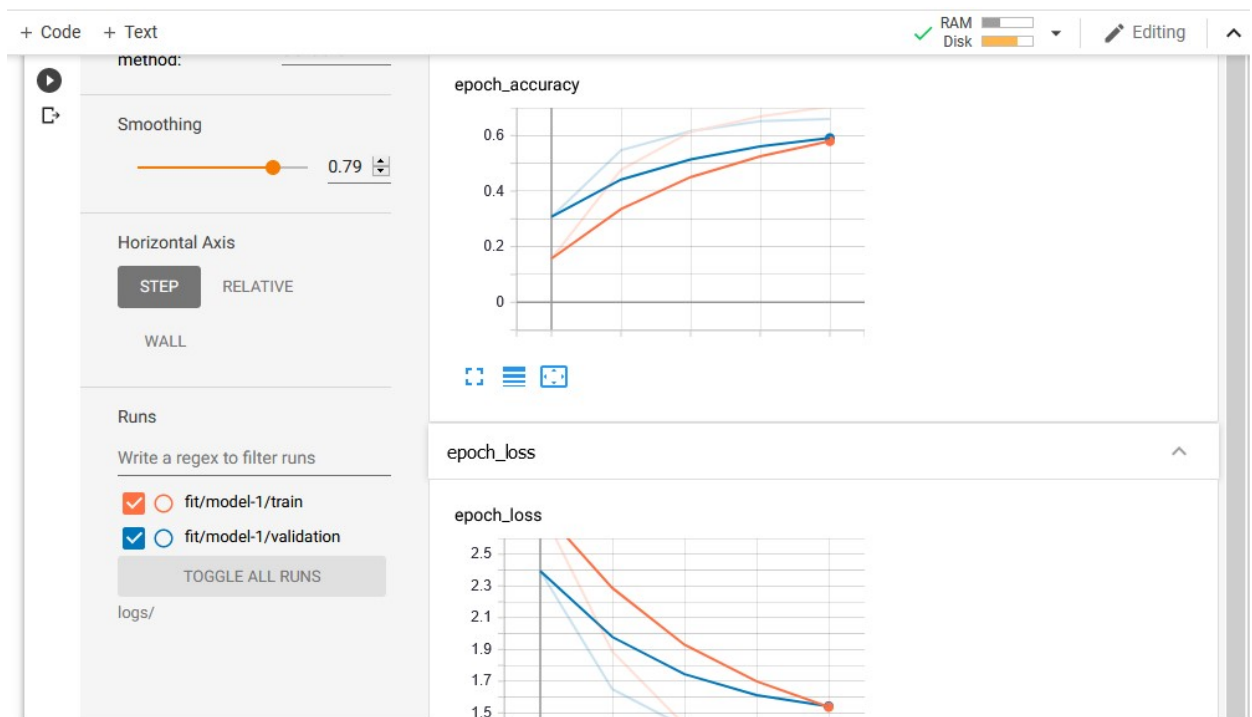
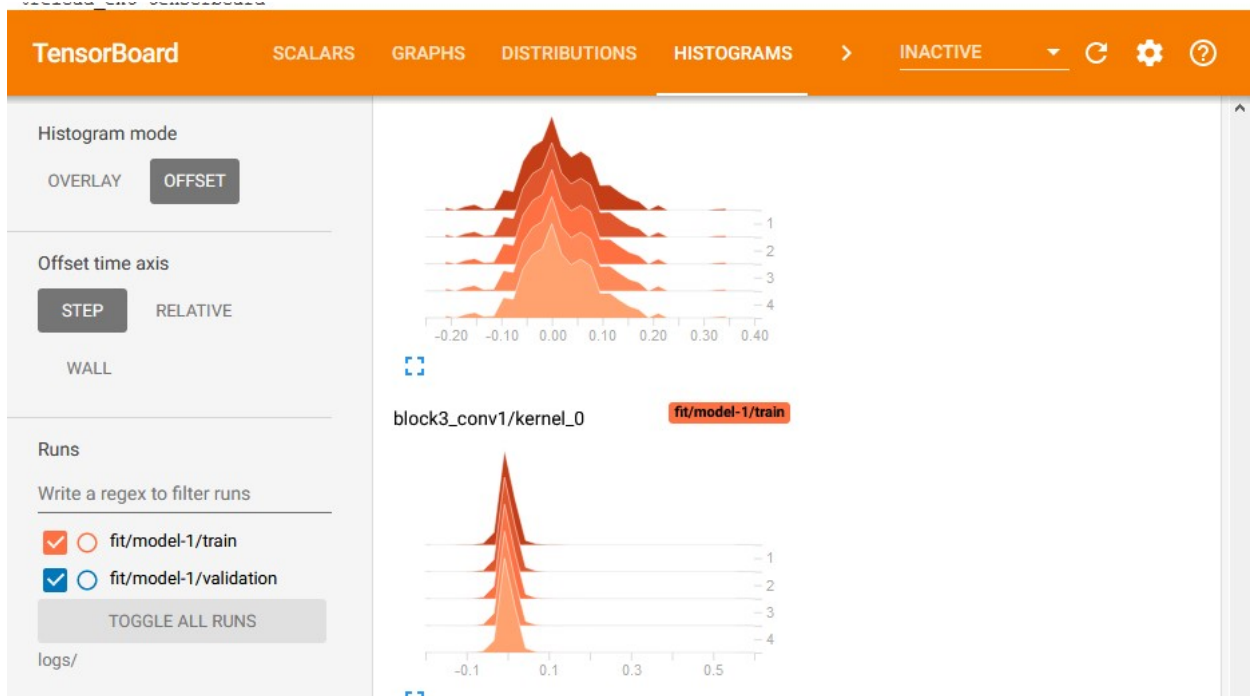
```
In [11]: %load_ext tensorboard
# Clear any logs from previous runs
#!rm -rf ./logs/
```

The tensorboard extension is already loaded. To reload it, use:

```
%reload_ext tensorboard
```

<IPython.core.display.Javascript object>





- 1. After each epoch model error is reducing and accuracy is increasing
- 2. Similary in the histogram we the kind of good distribution and see how they are changing from each epoch and layer to layer as well.
- 3. I got the more intuition on gram from here: <https://stackoverflow.com/questions/42315202/understanding-tensorboard-weight-histograms>

Model-2

1. Use [VGG-16](https://www.tensorflow.org/api_docs/python/tf/keras/applications/VGG16) (https://www.tensorflow.org/api_docs/python/tf/keras/applications/VGG16) pretrained network without Fully Connected layers and initialize all the weights with Imagenet trained weights.
2. After VGG-16 network without FC layers, don't use FC layers, use conv layers only as Fully connected layer. any FC layer can be converted to a CONV layer. This conversion will reduce the No of Trainable parameters in FC layers. For example, an FC layer with K=4096 that is looking at some input volume of size $7 \times 7 \times 512$ can be equivalently expressed as a CONV layer with $F=7, P=0, S=1, K=4096$. In other words, we are setting the filter size to be exactly the size of the input volume, and hence the output will simply be $1 \times 1 \times 4096$ since only a single depth column "fits" across the input volume, giving identical result as the initial FC layer. You can refer [this](http://cs231n.github.io/convolutional-networks/#convert) (<http://cs231n.github.io/convolutional-networks/#convert>) link to better understanding of using Conv layer in place of fully connected layers.
3. Final architecture will be VGG-16 without FC layers(without top), 2 Conv layers identical to FC layers, 1 output layer for 16 class classification. **INPUT --> VGG-16 without Top layers(FC) --> 2 Conv Layers identical to FC --> Output Layer**
3. Train only last 2 Conv layers identical to FC layers, 1 output layer. Don't train the VGG-16 network.

```
In [12]: from tensorflow.keras.applications.vgg16 import VGG16
#import tensorflow as tf
from tensorflow.keras.models import Model
from tensorflow.keras.optimizers import Adam

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dense
from tensorflow.keras.layers import Activation, Dropout, Flatten, Dense
import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator
tf.keras.backend.clear_session()
vgg16_model_2 = VGG16(input_shape=(224,224,3),weights = 'imagenet', include_top = False)
#vgg16_model_2.summary()

for layer in vgg16_model_2.layers:
    layer.trainable = False
    print(layer.name)
```

```

input_1
block1_conv1
block1_conv2
block1_pool
block2_conv1
block2_conv2
block2_pool
block3_conv1
block3_conv2
block3_conv3
block3_pool
block4_conv1
block4_conv2
block4_conv3
block4_pool
block5_conv1
block5_conv2
block5_conv3
block5_pool
Model: "vgg16"

```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0

```

In [13]: # use vgg16_model.output not vgg16_model_2.layers[-1].output
conv_layer_1 = Conv2D(4096, kernel_size=(7,7), strides=(1,1), padding='valid', kernel_initializer='he_normal')
conv_layer_2 = Conv2D(1000, kernel_size=(1,1), strides=(1,1), padding='valid', kernel_initializer='he_normal')
flatten      = Flatten()(conv_layer_2)
predictions = Dense(16, activation = 'softmax')(flatten)

```


```
model_2 = Model(inputs= vgg16_model_2.input, outputs = predictions)
model_2.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
conv2d (Conv2D)	(None, 1, 1, 4096)	102764544
conv2d_1 (Conv2D)	(None, 1, 1, 1000)	4097000
flatten (Flatten)	(None, 1000)	0
dense (Dense)	(None, 16)	16016
=====		
Total params: 121,592,248		
Trainable params: 106,877,560		
Non-trainable params: 14,714,688		
=====		

```
In [14]: from tensorflow.keras.utils import plot_model
```

Out[14]:



```
!rm -rf logs/
log_dir="logs/fit/model-2"
os.makedirs(log_dir, exist_ok=True)

tensorboard_callback =
tf.keras.callbacks.TensorBoard(log_dir=log_dir,histogram_freq=1)


model_2.compile(optimizer=Adam(learning_rate=0.001),loss='categorical_crossentropy',metrics=['accuracy'])
history = model_2.fit_generator(
    ImageGenerator_train,
    steps_per_epoch=38400//32,
    epochs=5,
    validation_data=ImageGenerator_validation,validation_steps= 9600//32,
    callbacks=[tensorboard_callback])
```

```
In [16]:
```

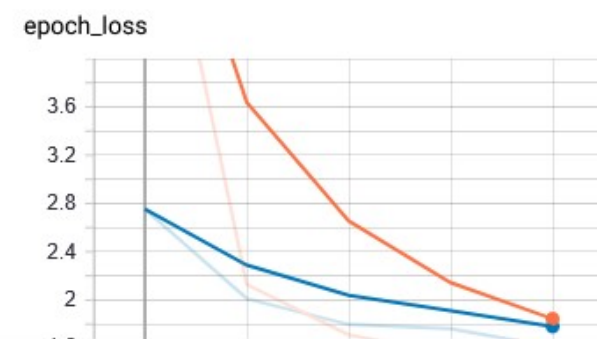
Reusing TensorBoard on port 6006 (pid 1159), started 0:45:49 ago. (Use '!kill 1159' to kill it.)

<IPython.core.display.Javascript object>

epoch_accuracy




epoch_loss




block1_conv2/bias_0

fit/model-2/train




block1_conv2/kernel_0


fit/model-2/train



block1_conv2/bias_0

fit/model-2/train





- 1. After changing INPUT --> VGG-16 without Top layers(FC) --> 2 Conv Layers identical to FC --> Output Layer each epoch model error is reducing and accuracy is trying to come closer
- 2. Similar in the histogram we see the kind of good distribution and see how they are changing from each epoch and layer to layer as well.
- 3. I got the more intuition on gram from here: <https://stackoverflow.com/questions/42315202/understanding-tensorboard-weight-histograms> (<https://stackoverflow.com/questions/42315202/understanding-tensorboard-weight-histograms>)
- 4. most of the weights are in the range of -0.15 to 0.15

Model-3

1. Use same network as Model-2 '**INPUT --> VGG-16 without Top layers(FC) --> 2 Conv Layers identical to FC --> Output Layer**' and train only Last 6 Layers of VGG-16 network, 2 Conv layers identical to FC layers, 1 output layer.

```
In [26]: from tensorflow.keras.applications.vgg16 import VGG16
#import tensorflow as tf
from tensorflow.keras.models import Model
from tensorflow.keras.optimizers import Adam

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Dense
from tensorflow.keras.layers import Activation, Dropout, Flatten, Dense
import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator
tf.keras.backend.clear_session()
vgg16_model_3 = VGG16(input_shape=(224,224,3),weights = 'imagenet', include_top = False)

number_of_last_layers_to_train = 6
non_trainable_layers_count = len(vgg16_model_3.layers)-number_of_last_layers_to_train
for layer in vgg16_model_3.layers[:non_trainable_layers_count]:
    layer.trainable = False
    print("Non-Trainable Layer:",layer.name)
for layer in vgg16_model_3.layers[non_trainable_layers_count:]:
    layer.trainable = True
    print("Trainable Layer:",layer.name)

vgg16_model_3.summary()
```

```

Non-Trainable Layer: input_1
Non-Trainable Layer: block1_conv1
Non-Trainable Layer: block1_conv2
Non-Trainable Layer: block1_pool
Non-Trainable Layer: block2_conv1
Non-Trainable Layer: block2_conv2
Non-Trainable Layer: block2_pool
Non-Trainable Layer: block3_conv1
Non-Trainable Layer: block3_conv2
Non-Trainable Layer: block3_conv3
Non-Trainable Layer: block3_pool
Non-Trainable Layer: block4_conv1
Non-Trainable Layer: block4_conv2
Trainable Layer: block4_conv3
Trainable Layer: block4_pool
Trainable Layer: block5_conv1
Trainable Layer: block5_conv2
Trainable Layer: block5_conv3
Trainable Layer: block5_pool
Model: "vgg16"

```

Layer (type)	Output Shape	Param #
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0

```

In [27]: # use vgg16_model.output not vgg16_model_2.layers[-1].output
conv_layer_1 = Conv2D(4096, kernel_size=(7,7), strides=(1,1), padding='valid', kernel

```

```
conv_layer_2 = Conv2D(1000, kernel_size=(1,1), strides=(1,1), padding='valid', kernel
flatten      = Flatten()(conv_layer_2)
predictions  = Dense(16, activation = 'softmax')(flatten)


model_3 = Model(inputs= vgg16_model_3.input, outputs = predictions)
model_3.summary()
```

Model: "model"

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
block1_conv1 (Conv2D)	(None, 224, 224, 64)	1792
block1_conv2 (Conv2D)	(None, 224, 224, 64)	36928
block1_pool (MaxPooling2D)	(None, 112, 112, 64)	0
block2_conv1 (Conv2D)	(None, 112, 112, 128)	73856
block2_conv2 (Conv2D)	(None, 112, 112, 128)	147584
block2_pool (MaxPooling2D)	(None, 56, 56, 128)	0
block3_conv1 (Conv2D)	(None, 56, 56, 256)	295168
block3_conv2 (Conv2D)	(None, 56, 56, 256)	590080
block3_conv3 (Conv2D)	(None, 56, 56, 256)	590080
block3_pool (MaxPooling2D)	(None, 28, 28, 256)	0
block4_conv1 (Conv2D)	(None, 28, 28, 512)	1180160
block4_conv2 (Conv2D)	(None, 28, 28, 512)	2359808
block4_conv3 (Conv2D)	(None, 28, 28, 512)	2359808
block4_pool (MaxPooling2D)	(None, 14, 14, 512)	0
block5_conv1 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv2 (Conv2D)	(None, 14, 14, 512)	2359808
block5_conv3 (Conv2D)	(None, 14, 14, 512)	2359808
block5_pool (MaxPooling2D)	(None, 7, 7, 512)	0
conv2d (Conv2D)	(None, 1, 1, 4096)	102764544
conv2d_1 (Conv2D)	(None, 1, 1, 1000)	4097000
flatten (Flatten)	(None, 1000)	0
dense (Dense)	(None, 16)	16016
=====		
Total params: 121,592,248		
Trainable params: 116,316,792		
Non-trainable params: 5,275,456		


```
In [28]: from tensorflow.keras.utils import plot_model
```

```
Out[28]:
```



```
In [29]:
```

```

!rm -rf logs/
log_dir="logs/fit/model-3"
os.makedirs(log_dir, exist_ok=True)

tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir, histogram_freq=

model_3.compile(optimizer=Adam(learning_rate=0.0001), loss='categorical_crossentropy',
history = model_3.fit_generator(
    ImageGenerator_train,
    steps_per_epoch=38400//64,
    epochs=3,
    validation_data=ImageGenerator_validation, validation_steps= 9600//64,
    callbacks=[tensorboard_callback])
  
```

```

Epoch 1/3
600/600 [=====] - 338s 563ms/step - loss: 6.5067 - accuracy: 0.4958 - val_loss: 4.9307 - val_accuracy: 0.6454
Epoch 2/3
600/600 [=====] - 343s 572ms/step - loss: 4.3560 - accuracy: 0.6787 - val_loss: 3.9051 - val_accuracy: 0.7113
Epoch 3/3
600/600 [=====] - 343s 571ms/step - loss: 3.4507 - accuracy: 0.7475 - val_loss: 3.2074 - val_accuracy: 0.7360
  
```

```
In [30]: %load_ext tensorboard
```

```

The tensorboard extension is already loaded. To reload it, use:
    %reload_ext tensorboard

<IPython.core.display.Javascript object>
  
```

```
In [ ]:
```

```

In [ ]: '''
        This code is referenced from the keras official documentation page for creating Image
        '''
  
```

```

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D
from tensorflow.keras.layers import Activation, Dropout, Flatten, Dense

from tensorflow.keras.preprocessing.image import ImageDataGenerator
batch_size = 64
# this is the augmentation configuration we will use for training
train_datagen = ImageDataGenerator(
    rescale=1./255,
    shear_range=0.2,
    zoom_range=0.2,
    horizontal_flip=True)

# this is the augmentation configuration we will use for testing:
# only rescaling
test_datagen = ImageDataGenerator(rescale=1./255)

tf.keras.backend.clear_session()

#import keras
#keras.backend.clear_session()
from tensorflow.keras.applications.vgg16 import VGG16
#import tensorflow as tf
from tensorflow.keras.models import Model
from tensorflow.keras.optimizers import Adam
#nb_epochs = 2
batch_size = 64
nb_classes = len(labels_dict)

vgg16_model = VGG16(input_shape=(224,224,3),weights = 'imagenet', include_top = False
x = vgg16_model.output

conv_layer_1 = Conv2D(4096, kernel_size=(7,7), strides=(1,1), padding='valid', activation='relu', input_shape=x.get_shape().as_list()[1:])
conv_layer_2 = Conv2D(1000, kernel_size=(1,1), strides=(1,1), padding='valid', activation='relu', input_shape=x.get_shape().as_list()[1:])
flatten = Flatten()(conv_layer_2)
predictions = Dense(16, activation = 'softmax')(flatten)


model_test = Model(inputs = vgg16_model.input, outputs = predictions)
for layers in vgg16_model.layers:
    layers.trainable = False
model_test.summary()


tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir,histogram_freq=100)

model_test.compile(optimizer=Adam(learning_rate=0.001),loss='categorical_crossentropy',metrics=['accuracy'])
history = model_test.fit_generator(
    ImageGenerator_train,
    steps_per_epoch=38400//32,
    epochs=5,
    validation_data=(test_datagen.flow_from_directory('data/test', target_size=(224, 224), batch_size=batch_size, shuffle=True)),
    validation_steps=1000//32,
    callbacks=[tensorboard_callback])

```


epoch_accuracy






block1_conv2/bias_0

fit/model-3/train



block1_conv2/kernel_0

fit/model-3/train



- 1.INPUT --> VGG-16 without Top layers(FC) --> 2 Conv Layers identical to FC --> Output Layer' and train only Last 6 Layers of VGG-16 network, 2 Conv layers identical to FC layers, 1 output layer - each epoch model error is reducing and accuracy is trying to come closer
- 2. Similary in the histogram we the kind of good distribution and seehow they are changing from each epoch and layer to layer as well.
- 3. I got the more intution on gram from here: <https://stackoverflow.com/questions/42315202/understanding-tensorboard-weight-histograms> (<https://stackoverflow.com/questions/42315202/understanding-tensorboard-weight-histograms>)
- 4. it is (mostly) equally likely for a weight to have any of these values, i.e. they are (almost) uniformly distributed, most of the weights are in the range of -0.15 to 0.15

